Agentic AI Agile Framework v 2.9 – A Comprehensive "People and Process-First" Playbook

Release Date: May 2025 **License: CC BY 4.0**

Author: Devashish Saxena (devashishsaxena@gmail.com)

Purpose: End-to-end operating model for conceiving, designing, testing, and governing enterprise-grade agentic AI systems. Assumes the prioritization of use cases has already been done as a separate exercise.
Audience: CDO / CIO, Product & Engineering Leaders, Transformation PMOs.

Phase 0: People and Process Discovery (Starting Point)

Purpose: Deeply understand the current state: who are the end users, who are the internal actors, what is the current process, how does it perform today, what works well, where are the friction points?

Activity	Description	Key Questions	Key Outputs
End user journey mapping	For customer facing use cases (e.g. customer support) define current end user journey and understand the different personas and especially where the current friction points are	Who are the end users? What is their journey today? What do they like about the current journey? Where are their friction points?	End user journey map by phases for different personas: what tasks are performed, where are current friction points, change-impact matrix
Current state Process Mapping	Build a common visual baseline of the current business process(es)	Current end-to-end flow? Bottlenecks, decision points, hand-offs?	Swim-lane map, pain-point heat-map
Business/Internal Stakeholder & Role Analysis	With the process map in place, capture who (internally) touches each step and their incentives/KPIs	Who does what? KPIs, incentives, friction?	RACI of employee actors, current pain points, change-impact matrix
Baseline Metrics Capture	With roles understood, pull in baseline hard numbers, understand trends	Current "impact" metrics revenue, NPS or outcome satisfaction and unit service cost	Baseline KPI dashboard (proof of impact), Data quality
Waste-to-Zero Workshop	Run a fast kaizen workshop – cross functional session	Which manual steps can be eliminated before automation?	Simplified future-state flow

	designed to identify and eliminate every non-value-added step		with “zero waste”, waste log
Knowledge Codification	In cleaned up process, identify the fastest, simplest error-free sequence of steps that achieves the desired business outcome	What is the “golden path” SOP for this workflow?	Canonical SOP deck for prompt/agent design, decision trees
Feature Opportunity Sizing	Size the steps for agentic lift (speed, experience quality, risk reduction) using chance-impact or impact-feasibility scoring	Where could autonomous agents lift speed, experience quality, risk?	Impact-feasibility matrix, prioritized use-case/feature backlog
Target-State Co-Design	With waste removed, SOPs codified, and opportunities ranked, design the future-state process that agents will inhabit.	How must the process evolve for autonomy & observability?	Future-state blueprint, re-engineered workflows

Outcome: Bundle all Phase 0 artifacts into a **Process-First Charter with baseline KPIs** that feeds Phase 1.

Trust-Posture Snapshot (read this first)

- Security grade: **ISO 27001 mapped**, zero hard-coded secrets.
- Privacy: **PII redacted at RAG retrieval**; row-level ACL.
- Kill-switch SLA: **< 30 s** tested quarterly.
- Model lifecycle: registry with upgrade checklist.

KPI Dictionary (enterprise-agnostic)

Metric	Definition	Why it matters
North-Star KPI	Single headline outcome (revenue, risk, experience)	Aligns agent design to business value

SSAT / NPS	Stakeholder-Satisfaction score (1-5) or Net Promoter Score	Proxy for adoption & quality
Autonomy %	Interactions fully handled by agent	Shows ROI realisation
Unit Service Cost	OPEX per completed interaction	Cost baseline & forecast
Escalation Rate	% routed to human oversight	Balance safety vs autonomy
Latency p95	95th percentile end-to-end time	Experience SLO
Policy Violations	Guard-rail breaches per 1k calls	Ethics & compliance health

Phase-Gate Calendar (typical 16-week pilot)

Wk 0-4 Phase 0 ——▶

Wk 5 Mission Definition Gate ✓

Wk 6-9 Phase 2 ——▶ Cost-to-Serve Gate ✓

Wk 10-12 Phase 3 ——▶ Ethics Gate ✓

Wk 13-14 Phase 4 ——▶ Prod Go / No-Go Gate ✓

Wk 15-16 Hyper-care roll-out

Phase 1. Mission Definition

Purpose: Turn the "Process-First Charter" from Phase 0 into a crystal-clear, metrics-anchored direction for the first set of agents to be built.

Key Activity	Description	Primary Roles	Key Outputs
Draft Agentic Epics	Convert each high-priority workflow into a single *Agentic Epic* statement: • Role (Sales-Assist Agent) • Goal (qualify and route inbound leads) • Tools/Data (CRM API, pricing DB) • Constraints (privacy tier, SLA)	AI Product Owner + Process Owner	Set of Epics—one per candidate agent

	• North-Star KPI (lead-conversion rate) • Optimization metric (cycle time)		
Define Success and Guard-Rails	• Quantify **North-Star KPI (revenue, cost, risk or customer/stakeholder experience) baseline vs target** • Select 2-4 supporting/operating KPIs (cost-per-unit, SSAT, error rate) • Establish guard rails - document policy, legal, ethical, safety, brand/tone and performance/cost constraints (e.g., no PII spill) • Specify escalation rules (route to human) if confidence thresholds not met	AI Product Owner, Ethics Partner, Risk Lead	**KPI & Guard-Rail Matrix** (one row per KPI, one row per guard-rail; includes target, owner, data source). **Escalation & Confidence Threshold Table** (links each trigger to the Responsibility Contract owner)
Responsibility Contracts	For each Epic assign: • **Agent Owner** (accountable exec) • Human On-Call (real-time override) • Failure Action (auto-pause, reroute)	Product Owner + Ops Lead	Updated Risk Register w/ contracts
Solution Architecture and Tech Feasibility Check	Align on high-level architecture (single agent vs multi-agent, RAG vs no-RAG, required tool integrations). Quick spike to confirm technical viability and token cost ballpark.	Agent Architect, Prompt Engineer, AgentOps Lead	Feasibility memo; rough infra sizing
Resource and Budget Alignment	Map required FTEs, sprint count, and infra spend. Ensure the **10-20-70** resource mix is still sensible (ensuring ongoing change/adoption activities)	Program PMO, CFO rep, Product Owner	Updated Cost-to-Serve model
Ethics and Alignment Pre-Check	Ethics Partner reviews Epics and guard-rails for bias, fairness, compliance. Flags items that must go through **Ethics Gate** later.	Ethics Partner	Pre-check sign-off or action items

Outcome: A formally approved **Mission Definition Deck** comprising of agentic epic 1-pagers, target KPIs, key guard-rails, responsibility contracts, solution architecture, technical feasibility, resource, and budget ballparks.

Phase 2. Agent Design & Tool Wiring

Purpose: Turn the approved "Mission Definition" from Phase 1 into a detailed, build-ready blueprint (prompts, memory design, data/tool wiring, security guard rails, and a validated cost-to-serve forecast).

Key Activity	Description	Primary Roles	Key Outputs
Platform and Buy-vs-Build Decision	Evaluate commercial / OSS agent frameworks (e.g., CrewAI, LangGraph, AutoGen) vs bespoke option. Select the stack that meets guard-rails, latency, extensibility, and TCO targets.	Agent Architect, AgentOps Lead, Security	• Platform decision note • Risk acceptance if bespoke
High-level Architecture and Memory Design	Choose cognition pattern (single agent, planner-executor, multi-agent). Define memory tiers (short-term token window, episodic DB, long-term vector DB, audit log) and planning loop/flow.	Agent Architect, Data Engineer	• Architecture diagram (planner, executor, memory tiers, tools, observability, security) • Memory schema • Planning loop spec (plan, act/execute, evaluate, record)
Tool and Data Integration Spec	List every external API, data product, or RAG corpus the agent will invoke. Document endpoints, auth, expected latency, cost limits, and observability hooks.	Prompt / Tooling Engineer, System SMEs	• Toolchain map • Security data-flow diagram
Prompt and Policy Engineering	Draft prompt taxonomy - system prompt, role/persona prompt, task prompt, function/tool wrappers, fallback prompts, tone guide,	Prompt Engineer, Ethics Partner	• Prompt library (version controlled)

	policy prompts (PII, ethics constraints). Include inline tags for confidence thresholds and escalation cues.		
Reusable Asset Library Contribution	Store new prompts, wrappers, eval configs in a shared Cross-Pod repository; tag with metadata for searchability.	Cross-Pod Guild delegate	Updated enterprise asset catalog
Security and Compliance Design	Threat-model the agent: auth scopes, rate limits, data classification, audit fields. Map to guard-rails and SOC2 / ISO / HIPAA controls as needed.	Security Architect, Ethics Partner	• Threat model matrix • Security requirements doc • Compliance mapping matrix • Ongoing security test plan
Evaluation Harness Set-up (a repeatable test case pipeline)	Build an automated test bed that objectively scores every new agent build against the KPIs and guard-rails defined in Phase 1—so failures are caught prior to production. Configure open harnesses (agentbench, AutoGen-eval, custom test suites) aligned to KPIs & guard-rails. Draft baseline scenarios.	Simulation/Test Engineer, AgentOps Lead	• Eval-config YAML / notebook
Prototype Spike and Cost Profiling	Build a thin vertical slice (happy path only) and run through evaluation harness to sample token, latency, and infra cost. Iteratively tune prompts / RAG chunking	Architect, Prompt Eng, Ops	• Cost-per-call range • Latency histogram
Cost-to-Serve Forecast and Stage-Gate Pack	Aggregate infra pricing, Ops FTE, 10-20-70 change mix. Verify data-quality readiness and produce "go / fix / defer" recommendation.	Product Owner, CFO rep, Ops Lead	• Cost Forecast model • Stage-gate deck

Phase 2 Critical Success Factors

1. **Prompt & tool wrappers checked into version control** with lint rules (no hard-coded API keys, no disallowed phrases).
2. **Security patterns baked-in early**—waiting until Phase 3 will create re-work.
3. **Cost-to-serve gate** signed off before heavy RLHF or large-scale data ingestion begins.
4. **Reusable assets pushed to the guild repo**—prevents each pod reinventing wrappers and eval configs.
5. **Evaluation harness covers the North-Star KPI and every guard-rail** (e.g., adverse prompt tests for policy).

Phase 3: Simulation & Safety Testing

Purpose: Validate agent behavior against functional KPIs and guard-rails in a fully sandboxed, risk-tiered environment before any end-user exposure.

Key Activity	Description	Primary Roles	Key Outputs
Simulation Environment Boot-up	Spin up sandbox infra, load snapshot RAG, install mocks; seed synthetic user IDs.	Test Eng, DevOps	Sandbox environment
Risk-Tiered Test Plan	Map each tool/data call to Tier 1/2/3; assign entry/exit gates	Test Eng, Security	Tiered test matrix
Synthetic and Edge-Case Dataset Build	Generate happy-path, edge, and stress datasets; include policy-violation probes.	Domain SME, Test Eng	tests/*.jsonl
Harness Execution and Metrics Capture	Run evaluation harness across all tiers; collect accuracy, policy, latency, cost.	AgentOps Lead	Raw run logs, metric CSV
Red-Team / Adversarial Blitz	Human red-teamers attempt jailbreak, PII extraction, cost abuse.	Red-Teamers, Ethics Partner	Red-team report, CVE list
Fallback-Path and Escalation Rehearsal	Force tool failures, low-confidence outputs; ensure escalation triggers fire.	Architect, Test Eng	Escalation drill report
Reinforcement Learning from Human Feedback	SMEs label 200–500 interaction pairs; tune model or prompt.	Prompt Eng, SME	Fine-tuned checkpoint / prompt v1.1

(RLHF) Micro-Sprint (optional)			
Safety Scorecard & Remediation Backlog	Consolidate results; tag blockers vs must-fix-later items.	Product Owner, Ethics Partner	Scorecard PDF; JIRA backlog
Ethics Gate Review	Present scorecard: sign-off, conditional go, or reject.	Ethics Board, Security, Product Owner	Formal Ethics approval

Outcome: Signed Ethics-Gate approval plus a Safety Scorecard showing accuracy, policy compliance, latency, and cost all within thresholds—clearing the way for limited human-feedback rollout.

Phase 4: Human Feedback & Iteration

Purpose: Expose the agent to real users in shadow or co-pilot mode, capture subjective trust signals, refine prompts/tools, and prove North-Star KPI lift without compromising safety.

Key Activity	Description	Primary Roles	Key Outputs
Shadow-Mode Launch	Agent runs in parallel to humans; outputs logged but not shown.	Ops Lead, Process Owner	Shadow log
Trust UX and Explainability Touchpoints	Inject confidence score, "why" button, tool call preview into UI.	Interaction Designer	Updated UI spec
User Education Bursts	5-min explainer videos, FAQ, slack posts.	Change-Enablement, Process Owner	Training artefacts
Weekly Adoption Huddle	Process Owner, Ops, Product review SSAT, override count, North-Star trajectory.	Change-Enablement, Process Owner	Huddle minutes, tweak list
Prompt / Tool Refinement	Apply tweaks from logs + huddle; bump prompt version.	Prompt Eng, Architect	Updated prompts file
KPI Delta Assessment	Compare live shadow KPIs vs baseline; update Cost-to-Serve forecast if needed.	Product Analyst	Delta sheet
Prod Go / No-Go Review	Steering committee checks KPI deltas, user-trust signals, open risks; decide.	Exec Sponsor, Product, Security	Signed Go / rollback plan

Success factors: SSAT ≥ baseline, override count trending down, trust cues understood, no unresolved Sev-1 issues.

Outcome: Production Go/No-Go decision backed by live SSAT, override, and cost data; updated prompt/tool version frozen for GA rollout.

Phase 5: Deployment, Monitoring & Drift Management

Purpose: Gradually roll out full autonomy, operate the agent under defined SLOs, and maintain performance through continuous drift detection, value realization reviews, and model lifecycle governance.

Key Activity	Description	Primary Roles	Key Outputs
Gradual Roll-Out Plan	5 % → 25 % → 50 % → 100 % traffic over "n" weeks with rollback checkpoints.	AgentOps Lead, Process Owner	Roll-out program plan
Observability Dashboard Go-Live	Build observability dashboard using e.g. Grafana/Datadog monitoring: latency, cost, autonomy score, policy violations.	Ops, DevOps	Live dashboard URL
Alert & SLO Configuration	Define p95 latency, cost per interaction, violation count SLOs; hook to incident management systems (e.g. PagerDuty/Opsgenie).	Ops, Security	Runbook & alert rules
Drift Detection and Re-alignment Loop	Weekly run: eval harness on fresh data measuring agent accuracy, cost and tone on fresh production logs flagging statistically significant degradation (compare to baseline; auto-ticket if KPI drop > accepted threshold)	Ops, ML Eng	Drift report; retrain tickets
Kill-Switch and Escalation Drills	Quarterly test of manual and auto shutdown; post-mortem.	Ops, Ethics Partner	Drill report

Regular Ongoing (e.g. Quarterly) Value-Realization Review	Baseline vs live KPI gap; ROI update.	Product Owner, CFO rep, Steering Committee	NorthStar KPIs actual vs target trend
Underlying Base Model Lifecycle Management	Governance and tooling to **version, monitor, upgrade, or deprecate** the underlying LLM or fine-tuned checkpoints.	Simulation/Test Engineer, AgentOps Lead	Model registry entries (e.g. MLflow)

Outcome: Agent in steady-state production with SLOs met, quarterly ROI verified, and active processes in place for drift re-alignment and future model upgrades.

Stage-Gates

1. Cost-to-Serve Forecast (after Design).
2. Ethics-Gate Approval (post Tier-2/3 Simulation).
3. Production Go / No-Go (post Feedback sprint).

RACI Heat-Map – Stage-Gates

Gate	Product	CISO	CFO	Process Owner	AgentOps Lead	Exec Sponsor
Mission Definition	A	C	I	R	I	A
Cost-to-Serve	A	C	A	R	C	I
Ethics Gate	C	A	I	R	C	A
Prod Go/No-Go	R	C	A	A	R	A

A=Approver, R=Responsible, C=Consult, I=Inform

References & Lineage

1. PwC (2024) *Agentic AI: The New Frontier*
2. McKinsey (2025) COO article on agentic pilots
3. BCG (2025) *AI Agents as the All-Stars*
4. AOSE literature (Wooldridge et al.)

5. OSS tool communities – LangChain, CrewAI, AutoGen, agentbench

Publication & Community Roadmap (appendix)

Steps 1-9 as outlined in prior guidance, including license, repo structure, CHANGELOG, first community call.

End of Playbook v 2.9